



Machine Learning Course


Week 2: Linear Regression

Ing 4 SI, 2022

Pr. Khadija SLIMANI

Course overview

1. Week 1 : Introduction to Data Science and Machine Learning
2. **Week 2: Univariate & Multivariate Linear Regression**
3. Week 3: Logistic Regression (Classification)
4. Week 4: Decision Trees (Regression & Classification)
5. Week 5: Model evaluation (overfitting, bias-variance, crossfolding, ...)



Course overview

1. Week 2 : Univariate & Multivariate Linear Regression
 1. Introduction to Linear Regression
 2. Simple Linear Regression
 3. Multiple Linear Regression
 4. Evaluation of a Linear Regression Model
 5. Practical Work



1.1

What is Linear Regression ?




What is Linear Regression ?

- ▶ **Linear Regression** is a statistical model used to predict the relationship between independent and dependent variables.
- ▶ In **regression** analysis, the **dependent variable** is denoted "Y" and the **independent** variables are denoted by "X".




Model Representation

- ▶ $X(i)$ denotes the “input” variables (house characteristics)
- ▶ $Y(i)$ denotes the “output” or target variable that we are trying to predict (price)
- ▶ A pair $(x(i), y(i))$ is called a training example
- ▶ A list of m training examples $(x(i), y(i)); i=1, \dots, m$ —is called a training set



$$h: \mathbf{X} \rightarrow \mathbf{Y}$$



Hypothesis or function that takes as input the house's characteristics to estimate its price.




1.2: Simple/Univariate Linear Regression



Simple Linear Regression: Model



Simple Linear Regression: Model Evaluation



Each line has different
parameters $\theta = (\theta_0, \theta_1)$
→ different errors

- Which line is the best fit ?
- **What should a good function $h_{\theta}(x)$ minimize?**
 - Sum error on all data points
 - Sum $\text{abs}(\text{error})$ on all data points
 - Sum error^2 on all data points



Simple Linear Regression: Model Evaluation

- ▶ Sum error on all data points.
- ▶ Sum $\text{abs}(\text{error})$ on all data points.
- ▶ Sum error^2 on all data points.



Sum error = 0



Sum $\text{abs}(\text{error}) > 0$




Sum $\text{error}^2 > 0$



Simple Linear Regression: Model Evaluation

- ▶ The Sum of Squared Errors (SSE) is also called Residual Sum of Squares (RSS)




$$SSE(\theta_0, \theta_1) =$$

$$\begin{aligned} & (y^{(1)} - (\theta_0 + \theta_1 * x^{(1)}))^2 + \\ & (y^{(2)} - (\theta_0 + \theta_1 * x^{(2)}))^2 + \\ & \dots\dots\dots + \\ & (y^{(m)} - (\theta_0 + \theta_1 * x^{(m)}))^2 \end{aligned}$$



Simple Linear Regression: Model Evaluation



- ▶ The green line is a better fit.
- ▶ Let's see how to find the best line automatically.



Simple Linear Regression: Cost function

- ▶ The cost function is the **technique of evaluating “the performance of our algorithm/model”**.
- ▶ It takes both predicted outputs by the model and actual outputs and calculates how much wrong the model was in its prediction.
- ▶ It outputs a higher number if our predictions differ a lot from the actual values.



Simple Linear Regression: Cost function

- ▶ The best hypothesis $h_{\theta}(x)$ is the one that minimizes the cost function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


Diagram illustrating the cost function $J(\theta_0, \theta_1)$ with annotations:

- Sum over all data points**: Points to the summation symbol $\sum_{i=1}^m$.
- Actual**: Points to $y^{(i)}$.
- Predicted**: Points to $h_{\theta}(x^{(i)})$.

$1/m$ - means we determine the average


$1/2m$, the 2 makes the math a bit easier, and doesn't change the weights θ we determine at all (i.e. half the smallest value is still the smallest value!)

- ▶ The learning algorithm should find $\theta^* = (\theta_0^*, \theta_1^*)$ that minimizes this cost
- ▶ Several algorithms:
 - ▶ Gradient descent
 - ▶ Ordinary least square (OLS): Used in the linear regression in python (sklearn)




Gradient Descent Algorithm: Intuition

- ▶ Let's assume that we have one parameter θ , and that we want to minimize $J(\theta)$: $h_{\theta}(x) = \theta x$
- ▶ GD starts by a random initial θ and iteratively update it to get towards θ^* .




- ▶ α is called the learning rate or the step size
 - ▶ **Too small:** Take baby steps -> Take too long to converge.
 - ▶ **Too large:** Can overshoot the minimum -> fail to converge.



Simple Linear Regression with GD

- ▶ Repeat until convergence :


$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$



- ▶ If we calculate the derivatives, the expression becomes:

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \quad \theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x^{(i)}$$

- ▶ Convergence means either:
 - ▶ The cost function is no longer changing by more than ϵ .
 - ▶ The number of iterations is reached



Ordinary Least Square Algorithm

- ▶ The Ordinary Least Square (OLS) approach chooses θ_0, θ_1 to minimize SSE.


$$SSE = \sum_{i=1}^m (y^{(i)} - \theta_0 - \theta_1 * x^{(i)})^2$$

- ▶ In this method, we will minimize SSE by explicitly taking its derivatives with respect to the θ_0, θ_1 , and setting them to zero.
- ▶ The minimizing values can be shown to be:

$$\hat{\theta}_1 = \frac{\sum_{i=1}^m (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2}$$



$$\hat{\theta}_0 = \bar{y} - \hat{\theta}_1 * \bar{x}$$

- ▶ Where $\bar{y} = \frac{1}{m} \sum_{i=1}^m y^{(i)}$, $\bar{x} = \frac{1}{m} \sum_{i=1}^m x^{(i)}$ are the sample means.

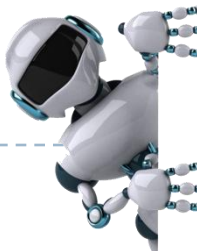


SSE is not perfect

- ▶ Which fit has larger SSE ?



- ▶ Larger SSE doesn't necessarily mean worst fit → Need an other metric.



R² Statistic

- ▶ R² Answers the questions: **“how much of variability in the output (y) is explained by the change in the input (x)”**.
- ▶ To calculate R², we use the formula:

$$R^2 = \frac{TSS - SSE}{TSS} = 1 - \frac{SSE}{TSS}$$
$$TSS = \sum_{i=1}^m (y^{(i)} - \bar{y})^2$$

- ▶ An R² statistic that is close to 1 indicates that a large proportion of the variability in the response has been explained by the regression.
- ▶ A number near 0 indicates that the regression did not explain much of the variability in the response.





1.3

Multivariate Linear Regression



Multiple Linear Regression

- ▶ In simple linear regression, we use one feature x to predict (y)
- ▶ In multiple linear regression, we have multiple features $X=(x_1, x_2, \dots, x_n)$

		X= input, features, covariate, predictors				y= output, target
		Size (m2)	Built Year	Nb bathrooms	Sale price (k\$)
m= number of training examples	$x^{(1)}$ 	200	2010	2	$y^{(1)} = 800$
	$x^{(2)}$ 	300	1995	2	$y^{(2)} = 750$

- ▶ $X_{(i)}$ is an n -dimensional feature vector
- ▶ $X_{(1)} = (200, 2010, 2, \dots)^T$ the feature vector of the first training example.
- ▶ $x_{j(i)}$ is the value of feature j in the i th training example. $x_{2(1)} = 2010$



Multiple Linear Regression

- In simple linear regression, $h_{\theta}(x) = \theta_0 + \theta_1 x$
- In multiple linear regression, $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$
- For convenience of notation, define $x_0 = 1$ ($x_0^{(i)} = 1$)


$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

- Rewrite in matrix notation

θ_0		x_0
θ_1		x_1
θ_2		x_2
...		...
θ_n		x_n

$\theta, x \in R^{n+1}$

$$h_{\theta}(x) = \sum_{i=1}^n \theta_i x_i = \theta^T x = \theta x^T$$



Multiple Linear Regression

- Cost function:


$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Simple Linear Regression, $n=1$

$$\begin{aligned} &\text{Repeat } \{ \\ &\theta_0 := \theta_0 - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})}_{\frac{\partial}{\partial \theta_0} J(\theta)} \\ &\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)} \\ &\quad \text{(simultaneously update } \theta_0, \theta_1 \text{)} \} \end{aligned}$$

Multiple Linear Regression, $n>1$

$$\begin{aligned} &\text{Repeat } \{ \\ &\quad \theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \\ &\quad \quad \text{(simultaneously update } \theta_j \text{ for } j = 0, \dots, n) \\ &\} \\ &\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)} \\ &\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)} \\ &\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)} \\ &\dots \end{aligned}$$



OLS for multiple features

- Cost function with matrix notations:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{2m} (X\theta - y)^T (X\theta - y)$$

- Where

$$X = \begin{bmatrix} - & (x^{(1)})^T & - \\ - & (x^{(2)})^T & - \\ & \vdots & \\ - & (x^{(m)})^T & - \end{bmatrix} \quad y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

- Derivatives with respect to θ $\nabla_{\theta} J(\theta) = X^T X \theta - X^T y$

- By setting them to zero $\theta = (X^T X)^{-1} X^T y$



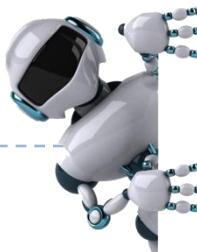
When to use OLS or GD ?

- ▶ The following is a comparison of GD and the OLS (normal equation):

Gradient Descent	OLS (Normal Equation)
Need to choose alpha	No need to choose alpha
Needs feature scaling	No need for feature scaling
Needs many iterations	No need to iterate
$O(kn^2)$	$O(n^3)$, need to calculate inverse of $X^T X$
Works well when n is large	Slow if n is very large

- ▶ $n=10^4$ - 10^5 is usually the threshold of choosing GD over OLS.
- ▶ Feature scaling helps converting the features to the same scale.
- ▶ For example, if x_i represents housing prices with a range of 100 to 2000 and a mean value of 1000, then,

$$x_i = \frac{\text{price} - 1000}{2000 - 100}$$



1.4

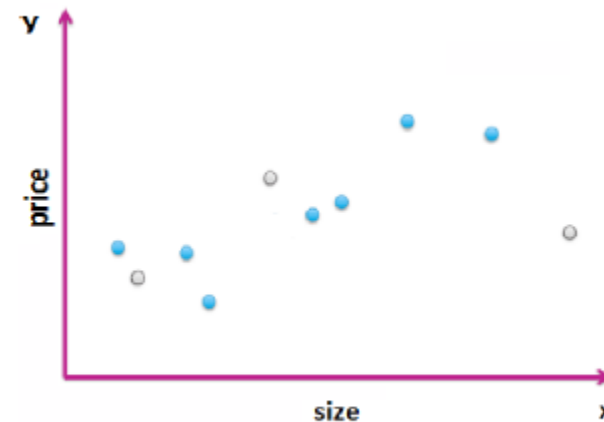
Assessing Performance



Assessing Performance

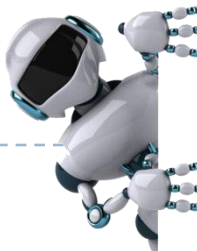
- ▶ This is about knowing how well the model will **generalize** to unseen data.
- ▶ One of the most used methods is splitting the data into train/test sets.

	Size	Price
Training set – 70%	2104	400
	1600	330
	2400	369
	1416	232
	3000	540
	1985	300
	1534	315
	1427	199
Test set 30%	1380	212
	1494	243



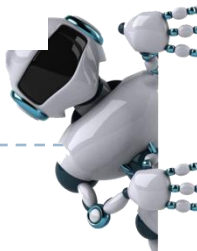


- ▶ Learn θ from training data (minimizing training error $J_{\text{train}}(\theta)$).

- ▶ Compute test error $J_{\text{test}}(\theta) = \frac{1}{2m_{\text{test}}} \sum_i^{m_{\text{test}}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$





Training error vs. model complexity

- ▶ Training error decreases with increasing model complexity





Training error vs. model complexity

- ▶ Test error can be used as an approximation of the generalization error




Bias – Variance tradeoff



- High bias model

The model does not capture enough
The structure of the training set
Parameters tend to be small


UNDERFIT



- High variance model

The model is too specific to the structure
Of the training set
Parameters tend to be very large

OVERFIT



1.5

Regularization



Regularization

- ▶ It's about finding balance between:
 - ▶ How well the model fits the data
 - ▶ The magnitude of coefficients
- ▶ This is achieved by incorporating a penalty on weights θ in the cost function.

- ▶ **Ridge Regression (L2 regularization)**

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=0}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2 \right]$$

- ▶ **Lasso Regression (L1 regularization)**

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=0}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j| \right]$$

- ▶ λ is the regularization parameter:
 - ▶ Ridge: Encourages small weights θ but not exactly 0.
 - ▶ Lasso: “Shrink” some weights θ exactly to 0.



GD with Regularization

- Reminder of the L2 regularization cost function:

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

- Previously:
$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$
$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$
- With regularization: $\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$
- α, λ are learning parameters to choose manually
- In practice: $(1 - \alpha\lambda/m)$ is between 0.99 and 0.95



Other Linear Regression Models

- ▶ Number of visiting customers to a website.
- ▶ Product demand, inventory, failure, ...
- ▶ Stock pricing.
- ▶ Insurance claims severity.

“Remember that all models are wrong;
the practical question is how wrong do they
have to be to not be useful.”

George Box, 1987



**Thank you for your
attention**



Practical work

LAB2: Back in 15min!

